

Lab 4 : Threads

September 25, 2018

Objective :

- Lab 4 is intended to give you some experience with programming using pthreads.

- **The Assignment: Threading Matrix Multiply**

In this assignment you are to write a version of matrix multiply that uses threads to divide up the work necessary to compute the product of two matrices. There are several ways to improve the performance using threads. You need to divide the product in row dimension among multiple threads in the computation. That is, if you are computing $A * B$ and A is a 10×10 matrix and B is a 10×10 matrix, your code should use threads to divide up the computation of the 10 rows of the product which is a 10×10 matrix. If you were to use 5 threads, then rows 0 and 1 of the product would be computed by thread 0, rows 2 and 3 would be computed by thread 1,...and rows 8 and 9 would be computed by thread 4.

- If the product matrix had 100 rows, then each "strip" of the matrix would have 20 rows for its thread to compute. This form of parallelization is called a "strip" decomposition of the matrix since the effect of partitioning one dimension only is to assign a "strip" of the matrix to each thread for computation. Note that the number of threads may not divide the number of rows evenly. For example, if the product matrix has 20 rows and you are using 3 threads, some threads will need to compute more rows than others. In a good strip decomposition, the "extra" rows are spread as evenly as possible among the threads. For example, with 20 rows and 3 threads, there are two "extra" rows ($20 \bmod 3$ is 2). A good solution will not give both of the extra rows to one thread but, instead, will assign 7 rows to one thread, 7 rows to another, and 6 to the last. Note that for this assignment you can pass all of the A and B matrix to each thread.

- What you need to do

Your program must conform to the following prototype:

```
my_matrix_multiply -a a_matrix_file.txt -b b_matrix_file.txt -t  
thread_count
```

where the `-a` and `-b` parameters specify input files containing matrices and `thread_count` is the number of threads to use in your strip decomposition.

The input matrix files are text files having the following format. The first line contains two integers: rows columns. Then each line is an element in row major order. Lines that begin with "#" should be considered comments and should be ignored. Here is an example matrix file

```
3 2
# Row 0
0.711306
0.890967
# Row 1
0.345199
0.380204
# Row 2
0.276921
0.026524
```

This matrix has 3 rows and 2 columns and contains comment lines showing row boundaries in the row-major ordering.

- Your assignment will be graded using a program that generates random matrices in this format (print-rand-matrix.c). You will need to write a function that can read matrix files in this format and to do the argument parsing necessary so that the prototype shown above works properly.
- Your program will need to print out the result of $A * B$ where A is contained in the file passed via the -a parameter and B is contained in the file passed via the -b parameter. It must print the product in the same format (with comments indicating rows) as the input matrix files: rows and columns on the first line, each element in row-major order on a separate line.
- Your solution will also be timed. If you have implemented the threading correctly you should expect to see quite a bit of speed-up when the machine you are using has multiple processors. For example, on some machine you may witness

```
my_matrix_multiply -a a1000.txt -b b1000.txt -t 1
```

completes in 8.8 seconds when a1000.txt and b1000.txt are both 1000 x 1000 matrices while applying

```
my_matrix_multiply -a a1000.txt -b b1000.txt -t 2
```

the time drops to 4.9 seconds (where both of these times come from the "real" number in the Linux time command).

Note this method of timing includes the time necessary to read in both A and B matrix files. For smaller products, this I/O time may dominate the total execution time. So that you can time the matrix multiply itself during development. A C function has been provided (c-timer.c) that

returns Linux epoch as a double (including fractions of a second) and it may be used to time different parts of your implementation. For example, using this timing code, the same execution as shown above completes in 7.9 seconds with one thread, and 4.1 seconds with two threads. Thus the I/O (and any argument parsing, etc.) takes approximately 0.8 seconds for the two input matrices.

- Hint: Refer to http://www.cs.ucsb.edu/~rich/class/cs170/labs/mmult/what_i_did.html

2. User-level threads

(a) Implementing user-level threads

In this section, you should implement a user-level thread library and a scheduler. To keep it simple, implement a round robin scheduler.

You will need to implement the following functions:

- `thread_create()`
- `thread_add_runqueue()`
- `thread_yield()`
- `thread_exit()`
- `schedule()`
- `dispatch()`
- `thread_start_threading()`

Each thread should be represented by a TCB (`struct thread` in the template) which contains at least a function pointer to the thread's function and an argument of type `void *`. The thread's function should take this `void *` as argument whenever it is executed. This struct should also contain a pointer to the thread's stack and two fields which store the current stack pointer and base pointer when it calls `yield`.

`thread create()` should take a function pointer and `void *arg` as arguments. It allocates a TCB and a new stack for the thread and sets default values. It is important that the initial stack pointer (set by this function) is at an address divisible by 8. The function returns the initialized structure.

`thread add runqueue()` adds an initialized TCB to the run queue. Since we implement a round robin scheduler, it is easiest if you maintain a ring of `struct thread`'s. This can be done by having a linked list where the last node points to the first node.

The static variable `current thread` always points to the currently executing thread.

`thread yield()` suspends the current thread by saving its context to the TCB and calling the scheduler and the dispatcher. If the process the resumed later, it continues executing from where it stopped.

`thread exit()` removes the caller from the ring, frees its stack and the TCB, sets `current thread` to the next thread to be executed, and calls `dispatch()`. It is important to dispatch the next thread right here before returning because we just removed the current thread.

`schedule()` decides which thread to run next. This is actually trivial because it is a round robin scheduler. Simply select the next thread in the ring. For convenience (e.g., for the dispatcher), it may be helpful to have another static variable which points to the last executed thread.

`dispatch()` actually executes a thread (the thread to run as decided by the scheduler). It has to save the stack pointer and the base pointer of the last thread to its TCB and restore the stack pointer and base pointer of the new thread. This involves some assembly code. In case the thread has never run before, it may have to do some initialization instead. If the thread's function returns, the thread has to be removed from the ring and the next one has to be dispatched. The easiest thing to do here is call `thread_exit()` since this function does that already.

`thread_start_threading()` initializes the threading by calling `schedule()` and `dispatch()`. This function should be called by your main function (after adding the first thread to the run queue). It should never return (at least as long as there are threads in your system).

In summary, to create and run a thread, you should follow the steps below:

```
static void
thread_function(void *arg)
{

    // Create threads here and add to the run queue if necessary

    while (condition) {

        do_work();
        thread_yield();
        if (condition)

            thread_exit();

    }
}

int
main(int argc, char **argv)
{
    struct thread *t = thread_create(f, NULL);
    thread_add_runqueue(t);
    // Create more threads and add to run queue if necessary
    thread_start_threading();
    printf("Done\n");
    return 0;
}
```